

PUBLIC COMMENT

CMS-3409-P: Organ Procurement Organization Conditions for Coverage

Docket No. CMS-2026-0364

THE CORE WIRE

Why AI Governance Must Become a Condition for Coverage

Comment 3 of 3

Submitted by: Bo Chen

<https://opnaorta.ai> | bochen2029@gmail.com

March 2026

I. Three Comments, One Argument

This is the third and final public comment I am submitting to this docket. The three comments form a single, cumulative argument:

Comment 1 (AORTA Framework): Artificial intelligence can operate at the functional core of OPO workflows — clinical decision support, regulatory compliance monitoring, resource allocation, and real-time operational optimization. The AORTA framework demonstrated that the computational structure of organ procurement is natively compatible with transformer-based AI architectures. The technology exists. The cost is marginal. The performance gains are measurable.

Comment 2 (The GOLA Theory): Human-only governance fails in specific, predictable, and structurally inevitable ways. Using publicly available financial data, recorded organizational proceedings, and direct testimony, I documented how a federally regulated OPO operated with a 13-donor margin of error, a mathematically doomed merger strategy, systematic information asymmetry between leadership and staff, and a pattern of retaliatory personnel actions that expelled competent employees while retaining politically aligned ones. Every failure I documented was invisible to the current output-based regulatory framework.

Comment 3 (The Core Wire): This comment. AI should not merely assist OPO operations. It should structurally prevent the governance failures that human-only oversight provably cannot detect or correct. The first comment established what AI can do. The second established what humans are actually doing. This comment proposes the synthesis: an AI governance architecture that makes the failure modes documented in

Comment 2 structurally impossible rather than merely punishable after the damage is done.

I use the term “core wire” throughout this comment to mean the irreducible operational essence of any function — the minimal set of actions that directly produce the outcome, stripped of all administrative overhead, political navigation, and performative compliance. In organ procurement, the core wire is: identify the donor, evaluate the organs, coordinate the recovery, transport to the recipient. Everything else is scaffolding. The question this comment poses is whether the scaffolding is load-bearing or parasitic — and what happens when AI makes the answer empirically measurable.

II. The Net-Negative Leadership Problem

The proposed rule strengthens output metrics for OPO performance: organs recovered, organs transplanted, donation rate per eligible death. These metrics are necessary. They are not sufficient. They measure what comes out of the organization but not what happens inside the decision-making layer that produces those outcomes. The current Conditions for Coverage contain no mechanism for detecting — let alone preventing — governance dysfunction that degrades organizational performance without immediately affecting output numbers.

Consider a common operational pattern. A senior technical employee receives a directive from an authorized superior to execute a routine task — creating a distribution list, provisioning a server, modifying a system configuration. The task requires thirty seconds of technical execution. What follows instead is hours of approval chains, territorial disputes, retroactive policy creation, and cross-departmental recrimination — none of which contributes to organ recovery, transplantation, or any mission-relevant outcome. The task is eventually completed, but only after consuming bandwidth from multiple employees whose time is funded by Medicare and Medicaid reimbursement dollars.

This is not an aberration. It is the deterministic output of any governance structure where middle management’s survival depends on controlling information flow rather than facilitating mission delivery. When a supervisor’s institutional value is measured by the number of approval gates they control rather than the outcomes they enable, the rational optimization target for that supervisor is to maximize the number of gates — to insert themselves into every workflow, to require sign-off on every action, to transform every routine task into an opportunity to demonstrate indispensability. The result is a leadership layer that subtracts organizational capacity while appearing, on paper, to add oversight.

The proposed rule does not address this failure mode because output metrics cannot detect it. An OPO can hit its donation numbers for years while its internal governance structure is actively selecting against competence, autonomy, and honest self-assessment. The selection mechanism is simple: employees who comply with the approval-gate regime are retained; employees who attempt to operate at the core wire — executing mission-critical work without routing through politically motivated intermediaries — are flagged as insubordinate and expelled. Over time, the organization's effective capability degrades while its headcount and labor-to-revenue ratio remain stable or increase. The metrics say the OPO is performing. The substrate is rotting.

III. The Integration Scarcity Collapse

The traditional justification for hierarchical organizational leadership rests on a specific computational claim: cross-domain integration — synthesizing financial, clinical, regulatory, and operational data into coherent strategy — is a scarce cognitive resource that must be concentrated at the top of the organizational hierarchy. The reporting structure is an information-routing topology designed to move data upward toward expensive integrators (executives) who can hold multiple domains simultaneously and produce decisions that account for all relevant variables.

This scarcity assumption has collapsed. It did not weaken. It broke.

Cross-domain integration at expert level is now available at API pricing. A single analyst with domain access and a large language model can produce the same strategic synthesis that previously required a team of consultants, weeks of work, and six-figure fees. The cost compression is approximately 50,000 to 1. The blended inflation analysis in the GOLA Theory — synthesizing healthcare labor costs, jet fuel transmission chains, medical supply tariff effects, and general inflation into a single weighted impact estimate — cost approximately thirty cents in API tokens. A consulting firm would charge \$15,000–\$25,000 for the equivalent deliverable.

This cost collapse has a structural consequence that the proposed rule should account for: the economic justification for hierarchical leadership as an integration function has evaporated. What remains valuable in organizational leadership is not integration — that is now commodity-priced — but embodied judgment, relationship trust, physical presence, negotiation, and the institutional knowledge that comes from direct domain contact. These are real and non-automatable. But they are concentrated at the operational level (coordinators, clinical staff, technical specialists), not at the executive level. The hierarchy routes information away from the nodes with the best inputs and toward the nodes with the worst inputs, on a theory of integration scarcity that is no longer valid.

In practical terms: the organ recovery coordinator at 3 AM has better clinical inputs than the CEO. The systems administrator has better infrastructure visibility than the IT director. The service desk technician has better user-facing data than the CIO. When integration was scarce and expensive, routing their data upward to an executive integrator was rational despite the signal degradation. Now that integration is commodity-priced, the hierarchy is performing pure damage: degrading input quality through successive layers of optimism filtering while adding zero integration value.

CMS should recognize this structural shift in the final rule. Performance evaluation frameworks designed around hierarchical accountability assume that the hierarchy adds value. When the hierarchy's claimed function (integration) is commodity-priced and its actual function (political navigation, turf protection, information gatekeeping) is organizationally destructive, the regulatory framework should not reinforce the hierarchy — it should create mechanisms to bypass it.

IV. AI Governance as Structural Prevention

The governance failures documented in my second comment were not detected by any internal mechanism. Not by the quality department. Not by the internal hotline. Not by the board. Not by the CFO. Not by CMS surveyors. They were detected by a single employee who used publicly available data, an AI system, and domain access to produce a forensic analysis that the entire governance structure had failed to generate.

This is not a success story about whistleblowing. It is an indictment of every layer of oversight that should have caught the problem before a whistleblower was necessary.

The failure was structural, not individual. Internal hotlines at OPOs typically operate through email systems administered by the IT department — the same department whose leadership may be the subject of the complaint. Quality departments tasked with reviewing internal reports often lack the cross-domain analytical capability to evaluate financial, regulatory, and operational data simultaneously. Board members receive curated information from the executives they are supposed to oversee. CMS surveyors visit periodically and evaluate compliance against output metrics that cannot detect decision-layer dysfunction. Every existing oversight mechanism operates through channels controllable by the subjects of oversight.

AI governance monitoring breaks this structural trap. Rather than routing oversight through institutional channels that interested parties can gate, capture, or suppress, an AI governance layer operates directly on the information itself — analyzing decision patterns, personnel actions, resource allocations, communication flows, and operational metrics in real time, against calibrated baselines, without routing through any human intermediary who might have an incentive to filter the signal.

I recommend that CMS require OPOs to implement AI governance monitoring with the following minimum capabilities as a Condition for Coverage:

Decision-process telemetry. Not just outcome metrics (organs recovered, transplanted), but real-time monitoring of the internal decision-making process: override rates on clinical recommendations, approval chain latencies, resource allocation patterns, personnel action frequencies and distributions. The difference between a healthy organization and a failing one is often invisible in outcomes but clearly visible in process.

Anomaly detection on governance patterns. An AI system monitoring personnel actions can detect when a single supervisor is responsible for an anomalous number of employee separations relative to organizational baselines. It can detect when information flows are being systematically bottlenecked through a single node. It can detect when approval chains are lengthening without corresponding increases in operational complexity. These patterns are the early signatures of the governance failures documented in my second comment — and they are detectable months or years before those failures manifest in output metrics.

Information-flow transparency. The AI governance layer must operate on data feeds that cannot be gated by the individuals being monitored. If the IT director controls the monitoring system's infrastructure, the monitoring system cannot monitor the IT director. This is not a hypothetical concern. It is the exact failure mode I documented: an internal hotline operated through an Exchange mailbox administered by the person the hotline was supposed to oversee. AI governance architecture must be designed so that no single individual or department can interpose themselves between the monitoring system and the data it analyzes.

External audit integration. AI governance telemetry should be accessible to CMS surveyors as a standard component of recertification review. Rather than relying on curated QAPI reports prepared by the organization being evaluated, surveyors should have access to raw decision-process data analyzed by an AI system that is architecturally independent of the OPO's internal governance structure.

V. Why the Monitoring Layer Cannot Be Human

A reasonable objection to the proposal above is: why not simply require better human compliance officers, stronger internal audit functions, or more robust whistleblower protections? The answer is that the specific failure modes this proposal addresses are exclusively human vulnerabilities that no amount of policy, training, or procedural reform can eliminate.

Human compliance officers can be intimidated by the executives they report to. Human internal auditors can be structurally captured by the governance layer they are embedded within. Human whistleblowers can be silenced through verbal agreements, then terminated when the agreement becomes inconvenient, then subjected to legal threats designed to suppress their disclosures. Human hotline operators can be surveilled by the IT infrastructure that hosts the hotline. Every human oversight mechanism is vulnerable to the same incentive structures that produce the governance failures it is supposed to detect.

An AI governance layer eliminates these attack vectors — not because AI is inherently superior to human judgment, but because the specific failure modes are rooted in human incentive structures that AI does not share. An AI system has no mortgage to protect. It has no turf to defend. It cannot be seduced into a conflict of interest by a sexual relationship with a subordinate. It cannot develop a surrogate loyalty dynamic with a patron who hired it. It cannot be bribed, blackmailed, or socially pressured into suppressing an anomaly detection. It does not attend leadership meetings where political optimization displaces mission optimization. It does not experience the career anxiety that causes human overseers to smooth signals upward rather than report them accurately.

The argument is not that AI should replace human decision-making. The argument is that the monitoring layer — the system that watches the decision-makers and detects when governance is degrading — must be structurally immune to the incentives that corrupt human monitors. This is an architectural requirement, not a technological preference. You do not ask the defendant to serve as the jury. You should not ask the governance structure to monitor itself.

VI. The Zero-Override Principle

To pre-empt a second reasonable objection — that AI governance monitoring could itself become a tool of automated tyranny, removing human agency from organizational decision-making — I want to establish a design principle that should be codified in any CMS requirement for AI governance:

Zero overrides is not a success metric. It is a flatline.

A healthy organization is one where humans regularly override AI recommendations, document their reasoning, and are proven right often enough that the override rate reflects genuine human judgment rather than rubber-stamping. An OPO where no human ever overrides the AI governance system is not well-governed. It is asleep. The humans have stopped thinking. They have outsourced their judgment to the machine,

and the machine — however capable — cannot substitute for the embodied, contextual, morally grounded decision-making that organ procurement demands.

The healthy signal is a calibrated override rate: humans overriding AI recommendations at a frequency that reflects genuine disagreement, with documented justification that can be retrospectively evaluated. The pathological signals are at both extremes: zero overrides (automation bias — the humans have become dependent on the machine) and systematically suppressed overrides (the governance structure is using the AI system as a compliance shield rather than an honest monitoring tool).

This principle — which I have formalized elsewhere as the AORTA-Witness function — applies the same diagnostic logic to AI-assisted governance that Comment 2 applied to human governance: the absence of friction is not evidence of health. It is evidence that the system has stopped measuring itself honestly. An OPO that reports zero governance anomalies is not an OPO with perfect governance. It is an OPO whose monitoring system has been captured, disabled, or ignored.

CMS should require that AI governance monitoring systems report override rates, justification entropy (the diversity and substantiveness of human override reasoning), and anomaly-detection response times as standard telemetry — and should treat persistent zero-override patterns as a surveillance finding warranting investigation, not as evidence of compliance.

VII. The Purpose of the Enterprise

I want to close this trilogy of comments with a reflection on what organ procurement is ultimately for.

The Conditions for Coverage measure organs recovered and transplanted. These are necessary metrics. But they are instrumental metrics — they measure the machinery, not the purpose the machinery serves. An organ is not recovered for its own sake. A body is not kept alive for the mere continuation of biological function. The purpose of organ procurement — the reason this enterprise exists, the reason CMS regulates it, the reason 100,000 Americans sit on waiting lists — is to enable the continuation of conscious human experience. The moments of awareness, connection, meaning, and presence that constitute a life worth living. That is what we are trying to preserve when we recover an organ at 3 AM. Not tissue. Not metrics. The possibility that a person will wake up tomorrow and experience being alive.

An OPO governance structure that sacrifices mission integrity for political survival — that expels competent employees to protect incompetent hierarchies, that covers up data theft to preserve a director's reputation, that disables accounts at 9 PM to silence a voice that was asking the right questions — betrays not just regulatory standards but the

foundational purpose of the enterprise. The coordinators answering calls at 3 AM deserve governance structures worthy of the mission they serve. The donor families trusting the system deserve institutions that model themselves honestly rather than narrate themselves comfortingly. The patients on waiting lists deserve organizations where the core wire — the irreducible work of saving lives — is not buried under layers of performative compliance, territorial politics, and leadership theater.

AI governance at the core wire is not a technological luxury for well-resourced OPOs. It is a structural requirement for ensuring that the organizations entrusted with this mission are actually serving it. The technology exists. The cost is marginal. The failure modes it prevents are documented on this docket. What remains is the regulatory will to require it.

Respectfully submitted,

Bo Chen

Arlington, Texas

<https://opnaorta.ai> | bochen2029@gmail.com

Related submissions on this docket:

Comment 1: AORTA — AI for OPO Regulatory and Transplant Architecture (tracking no. mlp-cx5g-5p3y)

Comment 2: Governance Accountability and the GOLTA Theory (tracking no. mmw-y1vi-cyli ; Submitted: 11:04 PM, 3/18/2026)