

*AORTA Research Program
Capstone Document*

The Same Shape

*Why Organ Procurement and Modern AI Architecture
Are Computationally Native to Each Other*

Bo Chen

March 2026

opnaorta.ai • MIT License

The magic of the fast inverse square root was not that Carmack found a clever approximation. It was that the answer was already inside the representation — encoded in the bit layout of IEEE 754 floating-point numbers — and the trick was recognizing that it was there.

— on substrate-native computation

1. The Question

In 1999, a piece of code appeared in the Quake III Arena graphics engine that has since become the most famous function in the history of software engineering. It computed the inverse square root of a floating-point number — an operation needed millions of times per frame for lighting calculations — using a technique so bizarre that its original source code included the comment “what the fuck?” beside the critical line.

The technique, now attributed to a lineage of programmers rather than to John Carmack alone, worked by exploiting a structural property of IEEE 754 floating-point representation: because the bit layout of a float encodes the number’s exponent in a position that overlaps with the integer interpretation of those same bits, treating the float’s bits as an integer and performing a single integer subtraction produces an approximation of the logarithm. One Newton’s method iteration refines it to usable accuracy. The result is three to four times faster than the standard library function.

The fast inverse square root is not interesting because it is fast. It is interesting because of *why* it is fast. It does not compute the answer more efficiently. It *reads* the answer from the substrate — from information that was already present in the data representation but that standard approaches ignored because they treated the representation as opaque. The float was not just storing a number. It was storing, for free, an approximation of that number’s logarithm. Carmack’s trick was recognizing the structural match between the representation’s properties and the computation’s requirements.

This document asks whether the Carmack insight can operate at a larger scale. Not one clever trick for one mathematical function, but a systematic structural match between an entire operational domain and an entire computational substrate — a case where the substrate’s native operations are the domain’s required operations, across every layer of the computational stack, so thoroughly that the architecture seems to design itself once you see the correspondence.

The domain is organ procurement. The substrate is GPU-accelerated transformer models. The claim is that they are the same shape.

2. The Domain: What Organ Procurement Computes

This section describes the computational structure of organ procurement operations without reference to AI. The goal is to characterize the domain on its own terms — what computations it requires, what makes them hard, and where human cognitive capacity is the bottleneck — so that when the substrate is introduced in the next section, the structural match can be evaluated against a description that was not reverse-engineered to fit.

2.1 The Core Wire

Every organ procurement organization (OPO) in the United States operates a process that begins with a hospital referral and ends, in successful cases, with organs transplanted into recipients. The process traverses a sequence of nodes: referral intake, donor evaluation, family approach and authorization, donor management, organ procurement, and allocation to recipients. This is the core wire — the end-to-end process that defines the organization's reason for existing. Everything else the organization does — hiring, training, compliance reporting, community outreach — exists to support the core wire.

Most organizations apply artificial intelligence to peripheral tasks: drafting emails, summarizing documents, generating reports. This is the computational equivalent of using the first digital computer to play tic-tac-toe. The ENIAC, the first general-purpose electronic computer, was built in 1945 for a specific purpose: computing artillery trajectory tables. The computation was expensive, the application was strategic, and the output — printed tables distributed to field units — was reusable. The ENIAC's designers did not waste the first available cycles of digital computation on convenience. They targeted the highest-leverage bottleneck in their operational domain.

The same principle applies. Fifty-six OPOs operate in the United States. The highest-leverage application of AI is not at the periphery of any individual OPO. It is at the core wire — and specifically, at the nodes where human cognitive labor is the bottleneck, where the computation is expensive and repeated, and where the output can be precomputed once and shared across the entire community.

2.2 Five Computations the Core Wire Requires

At each node of the core wire, the organization performs a specific type of computation. These computations differ in structure, dimensionality, and tractability, but they share a common property: all of them are currently performed by human cognitive labor under time pressure, cognitive load, and consequential stakes. A coordinator on a 24-hour shift integrates clinical data, applies regulatory policy, manages family dynamics, coordinates with transplant centers, and makes decisions that determine whether someone lives or dies — often while driving between hospitals at 3 AM.

The five computations, described in computational terms:

Computation 1: Authority Resolution. Given a donor's location, hospital type, organ type, and recipient characteristics, determine which regulatory authorities govern the case and what each authority requires. The inputs are discrete (jurisdiction, organ type, policy version). The operation is intersection — finding the overlapping requirements of multiple independent authority sets. The output is deterministic: for any given input combination, there is exactly one correct answer. The space of possible input combinations is large but finite and enumerable. This computation has the structure of a *compilation problem*: an expensive interpretive process (reading and reasoning across thousands of pages of policy) that can be performed once, offline, with the results stored in a lookup table for instant retrieval.

Computation 2: Clinical Viability Assessment. Given a donor's current labs, vitals, comorbidities, and temporal trends, assess whether each organ system remains viable for transplant and identify where the assessment is sensitive to small parameter changes. The inputs are continuous and high-dimensional (age, creatinine, creatinine trend, KDPI, BMI, vasopressor dependency, diabetes, hypertension, DCD status — for kidneys alone). The operation is evaluation of a complex function across a parameter space with interaction effects — every parameter modifies the significance of every other parameter. The output is not binary but graded, with uncertainty. This computation has the structure of a *phase space navigation problem*: determining the current state's position in a multi-dimensional landscape and its proximity to boundaries where the assessment changes qualitatively.

Computation 3: Deployment Compression. Whatever intelligence the system provides must run on a coordinator's laptop, in the field, without internet connectivity, without patient data leaving the device, and at a cost that 56 nonprofit organizations can sustain. The computation cannot be performed by a frontier model in the cloud at inference time. It must be performed by a frontier model offline, at compile time, with the results compressed into a model small enough to deploy locally. This has the structure of a *knowledge distillation problem*: transferring the capability of a large, expensive model into a small, deployable one through structured training data.

Computation 4: Allocation Matching. Given a set of available organs with specific quality profiles and a set of potential recipients with specific medical needs, compatibility constraints, and fairness requirements, find the assignment that maximizes transplant outcomes while respecting policy constraints. The inputs are two sets of high-dimensional vectors (donor organs and recipient profiles). The operation is constrained matching — finding optimal pairings that respect capacity limits (one organ per recipient, a finite number of organs per donor) and equity requirements (geographic priority, pediatric priority, sensitized patients). This has the structure of an *optimal transport problem*: finding the minimum-cost mapping between two probability distributions under constraints.

Computation 5: Relevance Communication. The coordinator does not need a dashboard. She needs a colleague who has already consulted every relevant data source and tells her, in natural language, what matters most right now. The inputs are the full clinical context (donor

data, map position, sensitivity gradients, temporal trends). The operation is context-dependent relevance determination — deciding which information is most important given this specific donor at this specific moment in this specific clinical context, and communicating it in natural language with appropriate emphasis. This has the structure of an *attention-mediated relevance reporting problem*: weighting inputs by context-dependent importance and generating a summary that preserves the relevance structure.

Table 1: The Five Computations of the Core Wire

Core Wire Node	Computation	Structure	Inputs	Key Property
Regulatory reasoning	Authority resolution	Compilation	Discrete, finite, enumerable	Deterministic; closed domain
Clinical assessment	Viability evaluation	Phase space navigation	Continuous, high-dimensional	Interaction effects; temporal dynamics
Field deployment	Capability compression	Knowledge distillation	Large model → small model	Offline compile, local runtime
Organ allocation	Constrained matching	Optimal transport	Two sets of vectors + constraints	Doubly-stochastic; fairness-aware
Coordinator interface	Relevance communication	Attention-mediated reporting	Full clinical context	Context-dependent weighting

These five computations were identified independently, over a period of months, by analyzing different nodes of the core wire. They were not selected to match any substrate. They were identified by asking, at each node: what is the coordinator actually computing when she makes this decision? The computational structure of each node was characterized before any architectural solution was proposed.

This matters for what follows.

3. The Substrate: What GPU-Transformer Systems Compute

This section describes the computational capabilities of modern AI infrastructure — specifically, transformer neural networks running on GPU hardware — without reference to organ procurement. The goal is the same as the previous section: characterize the substrate on its own terms, so the structural match can be evaluated honestly.

3.1 Steam and Steam Engines

The industrial revolution did not begin when someone discovered steam. Steam had been observed for millennia. The revolution began when someone built a steam engine — a mechanism that converted thermal energy into mechanical work directed at a specific task. The steam was potential. The engine was realization.

Modern AI has an abundance of steam. A B200 GPU — enterprise-grade, 192 gigabytes of high-bandwidth memory, capable of trillions of floating-point operations per second — can be rented for under three dollars an hour. Large language models with hundreds of billions of parameters are available through API calls at pennies per query. The raw computational power is extraordinary and historically cheap. What most organizations lack is not steam. It is the steam engine — the domain-specific architecture that converts raw computational power into operational value along their core wire.

The five native operations of the GPU-transformer substrate, described in computational terms:

Operation 1: Massively parallel independent evaluation. The GPU’s thousands of CUDA cores execute the same operation on thousands of different inputs simultaneously. This is the hardware’s defining capability. Any computation that can be expressed as “apply the same function to many independent inputs” runs at near-theoretical peak throughput. Perturbation testing, grid evaluation, batch inference, and pixel rendering are all instances of this operation class.

Operation 2: Self-attention (context-dependent relevance weighting). The transformer’s core mechanism computes pairwise relevance scores between all inputs in a sequence, normalizes them, and produces weighted representations. The operation is: given a set of elements, determine which elements are most relevant to which other elements in this specific context, and produce a representation that reflects that relevance structure. This was developed for natural language processing but is mathematically general — it operates on any structured set of inputs, not just words.

Operation 3: High-dimensional embedding and representation. The transformer maps inputs into a high-dimensional space (typically 4,096 or more dimensions) where geometric relationships encode semantic and structural relationships. Similarity is distance. Analogy is direction. Interaction effects that are invisible in the raw input space become linearly separable features in the embedding space. This is not a design choice but a mathematical consequence of

training: gradient descent finds representations where the relevant structure is geometrically accessible.

Operation 4: Knowledge distillation (capability compression). The standard technique for deploying large model capabilities on small hardware: a large “teacher” model generates training data, and a smaller “student” model learns to approximate the teacher’s function. Introduced by Hinton, Vinyals, and Dean in 2015, distillation has become a standard component of the ML engineering toolkit. The student doesn’t need to learn from scratch; it learns from the teacher’s outputs, compressing the teacher’s capability into fewer parameters.

Operation 5: Differentiable optimal transport (Sinkhorn attention). GPU-accelerated algorithms for solving optimal transport problems — finding minimum-cost matchings between distributions under constraints. The Sinkhorn algorithm iteratively normalizes rows and columns of a cost matrix to produce doubly-stochastic assignment matrices, and the entire operation runs as matrix multiplications on GPU. This was developed for machine learning applications (domain adaptation, generative modeling) but solves the general constrained matching problem.

Table 2: The Five Native Operations of the GPU-Transformer Substrate

Operation	Computational Primitive	Developed For	Mathematical Character
Parallel evaluation	Same function, many inputs, simultaneously	Graphics rendering	Embarrassingly parallel; $O(1)$ per additional input on GPU
Self-attention	Context-dependent pairwise relevance weighting	Natural language processing	Scaled dot-product; softmax normalization
High-dimensional embedding	Map inputs to space where structure is geometry	Representation learning	Learned nonlinear projection; distance = similarity
Knowledge distillation	Compress large model to small model via training data	Mobile deployment	Teacher-student; soft targets preserve distributional information
Sinkhorn optimal transport	Constrained minimum-cost matching between distributions	Domain adaptation	Doubly-stochastic matrix; iterative row-column normalization

These five operations were developed independently, across different research communities, over thirty years, for applications that have nothing to do with organ procurement. Graphics engineers built parallel evaluation for rendering pixels. NLP researchers built self-attention for machine translation. Representation learning researchers built embedding spaces for image classification. Deployment engineers built knowledge distillation for running models on phones. Applied mathematicians built Sinkhorn algorithms for computational geometry.

None of them were thinking about organ procurement. All of them were solving the *general form* of a problem that organ procurement happens to instantiate.

4. The Structural Match

Place the two tables side by side. The five computations the core wire requires (Table 1) and the five native operations the substrate provides (Table 2) are the same list, expressed in different vocabularies.

Table 3: The Structural Match

Core Wire Computation	Substrate Operation	Nature of Correspondence
Authority resolution (compilation)	Parallel evaluation on GPU	Both compile an expensive function into a reusable artifact through exhaustive offline evaluation
Clinical viability assessment (phase space navigation)	Self-attention + parallel perturbation testing	Both compute context-dependent relevance across multi-dimensional structured inputs and discover sensitivity boundaries
Deployment compression (field-deployable intelligence)	Knowledge distillation (teacher → student)	Both transfer capability from expensive models to cheap ones through structured training data
Allocation matching (constrained organ-recipient pairing)	Sinkhorn optimal transport	Both find minimum-cost assignments between two sets of vectors under capacity and fairness constraints
Relevance communication (natural language reporting)	Attention-mediated generation	Both determine which information is most relevant in context and produce a structured output reflecting that relevance

This is the central claim of this document: the structural correspondence in Table 3 is not a coincidence, not a metaphor, and not the product of reverse-engineering the domain description to fit the substrate. The five computations were identified by analyzing the core wire. The five operations were identified by cataloguing the substrate’s capabilities. They match because the computational requirements of organ procurement and the computational capabilities of GPU-transformer systems are, at the level of mathematical structure, *the same shape*.

We use the term **computational affinity** rather than isomorphism. Isomorphism would claim that the mathematical structure is identical in a formal, provable sense — that there exists a bijective mapping preserving all relevant structure. We do not make that claim. The domains differ in ways that matter: self-attention operates on learned continuous embeddings through differentiable operations, while clinical sensitivity detection operates on discretized parameter grids through Gaussian Process evaluation. The model’s learned relevance weights are statistical associations, not mechanistic causal relationships. Calling the correspondence an isomorphism would overstate the mathematical precision of the mapping.

Computational affinity means something precise: the computational primitives required by the domain are natively provided by the substrate, such that the substrate does not need to *emulate* the domain’s computations through an abstraction layer but can perform them *directly* using the

operations it was designed to execute. The distinction between emulation and native computation is the distinction between building a CPU inside Minecraft and running code on the actual CPU.

4.1 The Minecraft CPU Problem

The standard approach to applying AI in any domain is to treat the AI as a general-purpose reasoning engine: present the domain problem as a natural language prompt, let the model reason about it through token generation, and extract the answer from the generated text. This works. GPT-4 can reason about organ procurement cases if you give it enough context. Claude can analyze donor lab values if you paste them into a conversation. The outputs are often impressively sophisticated.

But this approach has the same structure as building a CPU inside Minecraft. Minecraft’s redstone circuits are Turing-complete — you can, in principle, build a functional computer inside the game. People have done it. The computer works. It executes instructions. It just does so at a millionth of the speed of the actual CPU running the game, because every operation passes through multiple layers of abstraction: the game engine interprets the redstone signal, which triggers a block update, which propagates through the circuit, which eventually flips a bit. The actual CPU performs the same logical operation in a single clock cycle.

When a coordinator pastes donor labs into ChatGPT and asks “should I pursue this kidney?,” the model reasons about the case through natural language token generation. Each token passes through the full transformer forward pass — billions of floating-point operations — to produce a single word. The model is using its computational substrate to *emulate* clinical reasoning through the medium of language, just as the Minecraft CPU uses redstone to emulate computation through the medium of a game engine. The emulation works. It is also enormously wasteful, because the substrate’s native operations (attention, embedding, parallel evaluation) are capable of performing the domain’s required computations directly, without the language intermediary.

The AORTA architecture is designed to run on the CPU, not inside Minecraft. The compilation pipeline uses GPU-parallel evaluation to build the lookup tables (native parallel operation, not language-mediated reasoning). The phase space map uses attention to detect clinical parameter sensitivity (native relevance weighting, not prompted reasoning about lab values). The deployment uses knowledge distillation to compress frontier capability into a local model (native training paradigm, not prompt engineering). The allocation uses Sinkhorn iterations on GPU (native matrix operation, not natural language matching). At every layer, the architecture uses the substrate’s operations *as they were designed to be used*, for the computations they were designed to perform.

This is what makes the structural match operationally significant rather than merely intellectually interesting. A Minecraft CPU and a real CPU both compute. The difference is six orders of magnitude in efficiency. The same class of difference separates a coordinator prompting ChatGPT about a donor case from a coordinator consulting a precomputed phase space map through a fine-tuned local model. Both produce an answer. One is native. The other is emulated.

4.2 What the Match Explains

If the structural match is real, it explains several properties of the AORTA architecture that would otherwise seem like fortunate engineering coincidences.

It explains why the compilation approach works for regulatory reasoning. If the regulatory domain is a finite, enumerable input space with deterministic outputs, and the GPU's native operation is exhaustive parallel evaluation of a function across input spaces, then *of course* the GPU can compile the regulatory domain into a lookup table. The hardware was designed for exactly this class of computation. The compilation is fast not because the engineers were clever but because the problem and the substrate are the same shape.

It explains why a transformer model fine-tuned on clinical data should develop sensitivity to clinical parameter interactions. The transformer's attention mechanism computes context-dependent relevance between inputs. Clinical viability assessment requires context-dependent relevance between parameters. If the model is given clinical parameters as inputs, its native attention operation will learn to weight parameters by their clinical relevance — not because it understands medicine but because attention is relevance weighting and the training data encodes which parameters are relevant to which outcomes. The substrate's native operation and the domain's required computation are, mathematically, the same operation applied to different inputs.

It explains why the dual-use trace format works for both context injection and fine-tuning. Knowledge distillation is a native training paradigm for transformers — the architecture was designed to learn from structured examples. A trace that teaches a model how to reason about a clinical case is performing distillation whether the trace is injected into the context window (in-context learning, which the transformer handles natively through attention over the prompt) or used as fine-tuning data (weight updates through backpropagation, which the transformer handles natively through gradient descent). The trace format's dual-use property is not a clever design choice. It is a consequence of the substrate's native capabilities.

It explains why Sinkhorn attention produces useful allocation matchings. The organ allocation problem is an optimal transport problem. Sinkhorn iterations solve optimal transport. The correspondence is not a metaphor — it is a mathematical identity. The domain's computation and the substrate's operation are literally the same algorithm.

And it explains why the entire architecture could be discovered incrementally, over months, through independent analysis of different core wire nodes, without advance planning. Each paper in the AORTA program was written to solve a specific problem at a specific node. The papers were not designed to converge. They converged because each node's computation matched a different native operation of the same substrate. The convergence was discovered, not designed, because the structural match is a property of the domain-substrate relationship, not of the engineering.

5. The ENIAC Principle: Pre-Compute Once, Deploy Everywhere

The structural match has a consequence that amplifies its leverage across the entire organ procurement community. If the domain's computations can be performed by the substrate's native operations, and those operations are expensive but deterministic (or approximately so), then they can be performed once, offline, and the results distributed as static artifacts.

The ENIAC computed ballistic trajectory tables once. The tables were printed and shipped to artillery units in the field. The field units did not need an ENIAC. They needed a book of tables and a person who could look up the right page. The expensive computation happened once, at a centralized facility with the necessary hardware. The deployment was cheap: paper, ink, distribution.

The AORTA framework applies the same principle at every node of the core wire. The regulatory lookup tables are compiled once on a rented GPU cluster and distributed to all 56 OPOs as a static artifact that any laptop can read. The phase space map is computed once, using Gaussian Process interpolation seeded by the ORCHID dataset (a peer-reviewed multi-center dataset of 133,101 donor referrals published in Nature Scientific Data) and refined by targeted frontier model evaluations, and distributed as a file. The fine-tuned narrator model is trained once on dual-use reasoning traces and distributed as a model checkpoint. Each artifact is expensive to create and cheap to deploy.

The community-level implications are profound. The United States has 56 OPOs. If each one independently builds its own AI infrastructure, the nation pays 56 times the development cost. If the artifacts are produced once and shared through an open framework — published under MIT license, documented in full, available at opnaorta.ai — the nation pays the development cost once and the deployment cost 56 times. The deployment cost is near zero: download the artifacts, install the model, begin using the system.

This is the ENIAC principle: **pre-compute the expensive thing once, distribute the cheap artifact everywhere.** It works because the structural match ensures that the domain's computations are compilable — they have a character that admits pre-computation. A domain whose computations were fundamentally unpre-computable (requiring real-time access to information that does not exist until inference time) would not benefit from the ENIAC principle. Organ procurement benefits because its core computations, while complex, operate over parameter spaces that can be characterized offline.

6. Where the Map Ends: The Human Line as Geometry

The structural match has a boundary, and the boundary is the document's most important contribution.

Not every computation along the core wire is native to the substrate. At the clinical viability assessment node, the coordinator integrates information that no data system can capture: the visual inspection of the organ during procurement, the subtle neurological signs that modify the clinical picture, the family dynamics that determine whether procurement can proceed on the necessary timeline, the clinical gestalt accumulated over years of practice that integrates signals below the threshold of explicit articulation. These are variables that exist only at the bedside, accessible only to the human.

In the AORTA program's earlier papers, the Human Line was presented as a policy boundary — a design decision about what AI should handle and what humans must decide. That framing was correct but incomplete. The structural match reveals a deeper characterization: the Human Line is a **topological boundary in information space** — the manifold where the substrate's representable dimensions end and the human's unmeasured dimensions begin.

The substrate can represent anything that can be encoded as numbers in a parameter space: lab values, vital signs, comorbidity flags, temporal trends, KDPI scores. These are the dimensions of the phase space map. The human accesses dimensions that are not in this space: the color of the kidney visible only during surgery, the surgeon's tactile assessment, the coordinator's accumulated experience with similar cases, the unspoken family dynamics that determine authorization timing. These dimensions are mathematically orthogonal to the map's axes — they contain information that is *linearly independent* of everything the system can represent.

This redefines the Human Line from a policy choice to a geometric fact. The system defers at the boundary not because someone decided it should defer, not because a more powerful model could overcome the limitation, but because the information required to resolve the case *does not exist* in the system's representable dimensions. The determination requires projecting into axes that only the human can access. No increase in model capability, no improvement in the map's resolution, no expansion of the parameter grid will ever cross this boundary, because the boundary is orthogonal to the substrate's entire representable space.

The output schema enforces this structurally. The AORTA narrator's response format contains fields for position, velocity, proximity to flip surfaces, sensitivity gradients, and uncertainty flags. It does not contain a field for determination. The absence is not an oversight. It is the topological boundary expressed as software architecture: the schema represents everything the substrate's dimensions can represent, and it represents *nothing* from the dimensions it cannot access. The coordinator reads the report. She has access to the orthogonal dimensions. She makes the determination. The Human Line is maintained not by policy enforcement but by information geometry.

This is what distinguishes the structural match from the standard AI-in-healthcare paradigm. Most clinical AI systems either ignore the boundary (producing determinations they are not informationally equipped to make) or impose the boundary externally (adding a disclaimer that says “please consult your physician”). The AORTA architecture derives the boundary from the geometry of the domain-substrate relationship itself. The boundary is discovered, not imposed. A system designed from the structural match *finds* the Human Line rather than having to be told where it is. Safety is not a constraint on the architecture. It is a property of the structural match.

7. The System That Learns to Need Itself Less

The structural match has a temporal property: the architecture improves through use, trending toward greater determinism as more cases are resolved and compiled into the static artifacts. This is not a feature added to the system. It is a consequence of the domain-substrate relationship.

When a coordinator at a border case overrides the system's assessment — proceeding where the map shows uncertainty or declining where it shows viability — the override is a new data point at specific coordinates in the phase space. The coordinates are known (the donor's parameters are in the system). The outcome is observable (the organ was or was not successfully transplanted). The override plus outcome becomes a new labeled anchor point that the Gaussian Process can incorporate, refining the map at exactly the coordinates where it was most uncertain.

Over time, the map's uncertainty decreases. The flip surfaces sharpen. The islands of stability expand. Regions that once triggered Tier 3 escalation (full human judgment required) accumulate enough labeled data to become Tier 2 (monitored navigation) or even Tier 1 (confident position report). The system's effective temperature decreases, like simulated annealing: hot at initialization (high uncertainty, many human overrides needed), progressively cooler as human knowledge is incorporated into the map.

This produces a specific organizational consequence: the cascade effect in reverse. Organizations traditionally scale complexity with scale — as volume grows, coordinators need more support staff, support staff needs managers, managers need analysts. The complexity of managing the complexity becomes its own burden. When the core wire's cognitive bottlenecks are addressed — not by replacing coordinators but by handling the systematic, compilable portion of their cognitive load — the support structures that existed to manage overwhelmed coordinators become unnecessary. Not because anyone is fired, but because the cognitive load that required the support never materializes. As coordinators leave through natural attrition, they are not replaced, because the system has absorbed the systematic portion of their work. The problem takes care of itself over time.

The system, in other words, learns to need itself less. As the map cools toward determinism, more cases fall into stability islands where a simple lookup suffices and the full narrator model's capabilities are unnecessary. The system's value migrates from “providing answers” to “identifying which cases need human judgment” — which is a more honest and more valuable function. The coordinator's cognitive capacity is freed for the cases that genuinely require the unmeasured dimensions only she can access.

8. Predictions: How to Test Whether the Match Is Real

A structural match that only explains retrospectively is not a discovery. It is a narrative. The test of whether the correspondence in Table 3 is real or imposed is whether it generates predictions that can be empirically confirmed or refuted. The structural match, if genuine, predicts at least six observable phenomena:

Prediction 1: Attention head specialization. If self-attention is natively performing clinical parameter sensitivity detection during fine-tuning, then the attention heads of a fine-tuned AORTA narrator model should show detectable specialization: specific heads should attend preferentially to the parameter pairs that the explicit phase space map identifies as high-sensitivity (e.g., creatinine + vasopressor status when the donor is near the renal flip surface). This is testable through standard attention probing techniques. If the attention patterns are diffuse and unstructured after fine-tuning on clinical data, the native computation claim is weakened.

Prediction 2: Explicit-implicit map agreement. The GP-interpolated phase space map and the fine-tuned model's internal geometry are two representations of the same viability function, computed through different mathematical processes. If the structural match is real, these representations should agree: querying both at the same parameter coordinates should produce correlated outputs. Divergences indicate either GP interpolation error or model hallucination, and either case is informative. This is testable through systematic paired evaluation across the parameter grid.

Prediction 3: Inter-OPO disagreement clusters at flip surfaces. If multiple OPOs contribute override data to a shared phase space map, the zones where OPOs disagree (one procures, another declines, at similar coordinates) should correspond to the high-uncertainty regions of the GP interpolation. Clinical disagreement IS flip surface proximity. This is partially testable now using ORCHID's multi-center data: comparing triage patterns across the six participating OPOs should reveal geographic clustering of divergence at specific parameter coordinates.

Prediction 4: The evaluation gate is cartographable. ORCHID shows that 85 percent of referrals — approximately 124,000 of 133,101 — are ruled out during initial evaluation. The triage decision is the first flip surface each referral encounters. The same phase space cartography methodology (discretize the parameter space, evaluate across the grid, discover sensitivity boundaries) should be applicable to the triage decision, revealing systematic divergence in how different OPOs draw the first flip surface. If the methodology works here, the structural match has generated a new architectural application from its own logic.

Prediction 5: Temporal trajectory prediction follows the Wisdom-Holman splitting. Donor management trajectories should decompose into a known baseline (physiological decline of a brain-dead donor without intervention) plus perturbations (effects of medical interventions: vasopressors, thyroid hormone, fluid management). If ORCHID's sequential event data is used to

train a trajectory model that separates baseline dynamics from intervention effects, and if the separated model outperforms a monolithic model that treats trajectories as undifferentiated sequences, the structural correspondence to Wisdom-Holman orbital splitting is empirically validated.

Prediction 6: Federated cooling converges faster than individual cooling. If 56 OPOs contribute override data to a shared phase space map — each override being a labeled point at specific coordinates, aggregated without sharing patient data — the national map’s uncertainty should decrease 56 times faster than any individual OPO’s map. This is the network effect prediction: the framework becomes more valuable as more organizations participate, and the improvement is quantifiable in the GP’s posterior uncertainty.

Table 4: Predictions and Their Status

Prediction	Tests	Status (March 2026)	If Confirmed	If Refuted
Attention head specialization	Probe fine-tuned model attention patterns	Untested (Phase 2 future work)	Native computation claim strengthened	Affinity framing weakened; match may be structural but not mechanistic
Explicit-implicit map agreement	Paired evaluation across parameter grid	Untested (requires compiled map)	Cross-validation architecture justified	Either GP or model is unreliable at specific coordinates
Inter-OPO disagreement at flip surfaces	ORCHID multi-center divergence analysis	Partially testable now	Flip surface concept empirically validated	Divergence may be driven by unmeasured dimensions, not parameter space
Evaluation gate cartography	Phase space methodology applied to triage decision	Untested (separate research program)	Structural match generates new applications	Methodology may not generalize from assessment to triage
Wisdom-Holman trajectory splitting	Separate baseline from interventions in ORCHID temporal data	Untested (requires temporal modeling)	Symplectic principle extends from metaphor to method	Clinical trajectories may not decompose cleanly
Federated cooling convergence	Multi-OPO override aggregation	Untested (requires multi-OPO deployment)	Network effect quantified	Override data may be too noisy or biased for aggregation

We are explicit: as of March 2026, none of these predictions have been empirically tested. The AORTA program has produced architectural designs, evaluation protocols, and a detailed research agenda. It has not produced experimental results. The predictions are stated here so that they can be tested, and the document’s claims can be evaluated against outcomes rather than

against rhetoric. A structural match that survives empirical testing is a discovery. A structural match that fails empirical testing was a plausible but incorrect hypothesis, honestly stated. Either outcome advances the field.

9. What This Document Does Not Claim

The structural match is specific. This document claims that one operational domain (organ procurement) and one computational substrate (GPU-transformer systems) exhibit computational affinity across five identified layers. It does not claim that this match is universal.

Other domains may exhibit the same kind of structural match with the same or different substrates. Emergency response coordination, air traffic control, supply chain logistics, and intensive care unit management all involve structured decision-making under uncertainty with context-dependent relevance across high-dimensional parameter spaces and temporal dynamics. Whether the structural match extends to these domains is an empirical question that this document does not answer. It would be irresponsible to claim universality from a single domain-substrate pairing, regardless of how thorough the match appears.

This document also does not claim that the AORTA architecture has been validated. The architecture exists as a design, described across six published papers and implemented in an open framework. The evaluation protocols are defined but not executed. The phase space map has been specified but not compiled. The narrator model has been selected but not fine-tuned. The voice interface has been designed but not benchmarked. The six predictions in Section 8 are untested. The document makes architectural claims, not experimental claims, and it is transparent about the distinction.

Finally, this document does not claim that the structural match eliminates the need for human clinical judgment. The match has a boundary — the Human Line — and that boundary is a geometric fact, not a temporary limitation. The dimensions accessible only to the human at the bedside (visual organ assessment, clinical gestalt, family dynamics, surgical judgment) are orthogonal to the substrate's representable dimensions. No expansion of the substrate will absorb these dimensions, because they are not the kind of information that digitizes. The system maps the terrain the substrate can represent. The human navigates the dimensions the substrate cannot. The Human Line persists because it is real.

10. The Papers: Six Architectural Contributions

Each paper in the AORTA program addresses a specific computation at a specific core wire node. They were written independently, over a period of months, without advance planning for convergence. This section provides a one-paragraph summary of each for readers encountering the program for the first time. Full details are published at opnaorta.ai under MIT license.

Reasoning Trace Injection (RTI) establishes the methodology for transferring domain-specific reasoning competence to any language model. The core innovation is a dual-use trace format: the same JSONL dataset works for context injection (teaching a closed model like Claude or GPT-4 the OPO reasoning pattern within its context window) and for fine-tuning (permanently embedding the reasoning pattern in the weights of an open model like Qwen3.5-9B). Layer A traces are public and domain-general; Layer B traces are organization-specific. The dual-use property is the distillation mechanism that enables the ENIAC principle: generate the expensive traces once, deploy everywhere.

Reasoning as Infrastructure introduces the safety architecture: a formal grammar that constrains the model's output space, a compliance automaton that monitors for Human Line violations in real time, and the `BOUNDARY_ENFORCE` signal that triggers automatic escalation when the system approaches the orthogonal boundary. The safety properties are structural, not behavioral — they are enforced by the grammar and the automaton, not by the model's training.

The Bimodal Core designs the unified dataset architecture that separates deterministic facts (which compile into Tier 1 lookups) from probabilistic assessments (which require Tier 2 or Tier 3 reasoning). The bimodal structure ensures that deterministic knowledge is never re-derived at inference time and probabilistic knowledge is always flagged with calibrated uncertainty.

Native Intelligence introduces substrate-native computation: the argument that GPU hardware and transformer architectures should be used for their native operations rather than as emulation layers for natural language reasoning. The paper applies Sinkhorn attention to donor-recipient allocation, demonstrating that the allocation problem is structurally identical to optimal transport — a native GPU operation. The Carmack insight (reading the answer from the substrate rather than computing it through abstraction) is introduced here.

The Compiled Core proves that regulatory reasoning in organ procurement can be compiled into deterministic lookup tables. A frontier model serves as the compiler, evaluating every intersection of governing authorities offline and producing static lookup infrastructure that a coordinator's laptop queries in milliseconds. The paper introduces the three-tier runtime (Tier 1: deterministic lookup, Tier 2: grammar-guided generation, Tier 3: full reasoning with human escalation) and the self-obsoleting property (the system trends toward Tier 1 through use).

The Flip Surface extends the compilation methodology from regulatory reasoning (fully compilable) to clinical viability assessment (partially mappable). Phase space cartography discretizes the clinical parameter space, evaluates viability across the grid using a three-tier

evaluation function (deterministic rules, ORCHID-derived anchor points, frontier model oracle), discovers flip surfaces through GPU-parallel perturbation testing, and delivers the map through a fine-tuned narrator model that the coordinator can consult by voice. The paper redefines the Human Line as a topological boundary in information space and introduces the cooling mechanism.

11. The Coordinator

It is 3:14 in the morning. The coordinator's phone has rung. A 67-year-old male with a massive intracerebral hemorrhage has been declared brain dead. The referral has been initiated. The coordinator is driving to the hospital. She has been on shift for fourteen hours. She will be on shift for ten more.

She does not know that the regulatory lookup table she will query in thirty seconds was compiled by a frontier model that evaluated every authority intersection in the OPTN policy manual, running 86,000 parallel evaluations on a GPU that was rented for three dollars an hour — and that the same table is running simultaneously on coordinator laptops at all 56 OPOs in the nation, compiled once and distributed everywhere, like artillery trajectory tables printed and shipped from a computing facility in 1945.

She does not know that the phase space map that will tell her where this donor's kidneys stand — creatinine rising, KDPI at 78, approaching a sensitivity boundary in roughly four hours — was built by seeding a Gaussian Process with labeled outcomes from 133,101 referrals across six OPOs, refined by 200 targeted frontier evaluations at the points of highest uncertainty, and stress-tested by a GPU running 1.5 million perturbation evaluations in parallel to discover the exact boundaries where the clinical assessment reverses.

She does not know that the model answering her voice query — “How are the kidneys looking?” — is a 9-billion-parameter transformer running on her laptop's GPU, fine-tuned on reasoning traces that were designed to work both as in-context examples and as training data, distilled from a model two orders of magnitude larger, and that the attention heads in its forward pass are natively computing the context-dependent relevance between this donor's parameters because that is what attention does, because the mathematical operation that transformed machine translation also transforms clinical sensitivity detection, because the problem and the substrate are the same shape.

She does not know that the system will not tell her whether to pursue the kidney. It will tell her where she is, how fast things are changing, and where the boundaries are. It will not cross the Human Line, because the Human Line is not a policy the engineers imposed but a geometric fact the architecture discovered — the boundary where the measurable dimensions end and the dimensions only she can access begin. The color of the kidney. The surgeon's assessment. The family's readiness. Her own judgment, honed by years of standing in operating rooms at 4 AM, integrating signals that no parameter grid will ever capture.

She does not need to know any of this. She needs to know that the kidneys are in a watchable zone, that the creatinine trend is the axis to monitor, that she should order a repeat BMP in two hours, and that the liver and heart are stable. She needs a colleague who has already done the homework and can tell her what matters right now.

Five research communities — graphics engineers, NLP researchers, ML deployment engineers, computational geometers, applied mathematicians — spent thirty years building the components she needs, without knowing they were building them for her. They were solving the general forms of the problems she faces every night. The attention mechanism, the knowledge distillation framework, the GPU parallel evaluation architecture, the optimal transport algorithm, the Gaussian Process interpolation — all developed for other purposes, all converging on her 3 AM phone call because the computational structure of her work and the computational structure of their inventions are, at the deepest level, the same.

The AORTA framework is the steam engine that channels their steam into her core wire. The papers describe how each component works. This document describes why they all work together: because the problem and the substrate are the same shape, and when you discover a structural match that thorough, the architecture designs itself.

The framework is released under MIT license at opnaorta.ai. The methodology, the datasets, the evaluation protocols, the model specifications, and the compiled artifacts are published in full. The infrastructure for organ procurement should be as open as the mission it serves.

— — —

Every organ saved is a life continued.

References

- [1] Chen, B. (2026). Reasoning Trace Injection: Transferring Domain-Specific Reasoning Competence to Any LLM via Context Window. AORTA Research Program. opnaorta.ai.
- [2] Chen, B. (2026). Reasoning as Infrastructure: Grammar-Guided Safety Architecture for AI-Assisted Organ Procurement. AORTA Research Program. opnaorta.ai.
- [3] Chen, B. (2026). The Bimodal Core: A Unified Dataset Architecture for Deterministic and Probabilistic Organ Procurement Reasoning. AORTA Research Program. opnaorta.ai.
- [4] Chen, B. (2026). Native Intelligence: Substrate-Native Computation and Geometric Allocation in Organ Procurement. AORTA Research Program. opnaorta.ai.
- [5] Chen, B. (2026). The Compiled Core: A Three-Tier Deterministic Architecture for Regulatory Reasoning in Organ Procurement. AORTA Research Program. opnaorta.ai.
- [6] Chen, B. (2026). The Flip Surface: Phase Space Cartography for Clinical Decision Support in Organ Procurement. AORTA Research Program. opnaorta.ai.
- [7] Adam, H., Pollard, T., Suriyakumar, V., et al. (2026). Organ retrieval and collection of health information for donation: The ORCHID dataset. *Scientific Data*, 13, 120.
- [8] Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the Knowledge in a Neural Network. arXiv:1503.02531.
- [9] Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems*, 30.
- [10] Cuturi, M. (2013). Sinkhorn Distances: Lightspeed Computation of Optimal Transport. *Advances in Neural Information Processing Systems*, 26.
- [11] Rasmussen, C.E. & Williams, C.K.I. (2006). *Gaussian Processes for Machine Learning*. MIT Press.
- [12] Wisdom, J. & Holman, M. (1991). Symplectic maps for the N-body problem. *The Astronomical Journal*, 102, 1528–1538.
- [13] Lomonaco, C. (attributed). Fast inverse square root algorithm. id Software, Quake III Arena source code (1999).
- [14] Qwen Team. (2026). Qwen3.5 Technical Report. <https://qwen.ai>.
- [15] Organ Procurement and Transplantation Network (OPTN). Policy documentation. <https://optn.transplant.hrsa.gov/policies-bylaws/>.
- [16] AORTA Framework. (2026). Autonomous OPO Reasoning and Triage Architecture. <https://opnaorta.ai>.